



# EchoDraft — AI Builders Challenge Project Spec (July 2026)



**EchoDraft** — a voice-preserving content repurposer. Paste one pillar piece → get platform-native versions that keep *your* voice. Built for the IBM AI Builders Challenge, July theme: *Reimagine Creative Industries with AI*. **Deadline: July 31, 2026, 11:59 PM ET (~Aug 1, 5:59 AM SAST).**

## 1. The problem

Two of the loudest 2026 creator complaints, in one product:

- **Repurposing rarely works.** Creators are told to "post everywhere," but the same message copy-pasted across platforms flops — hooks, pacing, and tone need real adaptation per platform.
- **AI flattens your voice.** Writers fear AI makes everything sound like "an amalgamation of the entire internet." 2026 audiences reward authentic, on-brand content over polish, so generic AI output actively hurts creators.
- **Disconnected tools = manual labor.** Creators act as the human bridge, copy-pasting between research, writing, and publishing tools that don't share context.

## 2. The solution

EchoDraft captures a creator's **voice profile** from a few samples, then transforms a single pillar piece into multiple **platform-native** posts — each with its own hook and pacing, all rewritten in the creator's own voice.

| One input. Many platforms. Still sounds like you.

### 3. Target users

- Solo content creators and "solopreneurs" juggling multiple platforms
- Small business / lean marketing teams repurposing blog posts
- Freelance writers and coaches building a personal brand

### 4. Core features (MVP)

1. **Voice Profile Builder** — user pastes 3–5 samples of their own writing; AI extracts tone, vocabulary, sentence rhythm, and "do/don't" traits into a reusable profile.
2. **Pillar Input** — paste a blog post, transcript, or long caption (the source content).
3. **Platform-native generation** — one click produces versions for:
  - LinkedIn post
  - X / Twitter thread
  - Instagram caption
  - YouTube description
  - Email newsletter blurb

Each uses a platform-specific hook, length, and pacing — not a copy-paste.

4. **Voice-match rewrite** — every output is conditioned on the saved voice profile so it reads like the creator, not generic AI.
5. **Copy / export** — one-click copy per platform, plus "export all."

### Nice-to-have (if time allows)

- **Authenticity check** — flags generic "AI-sounding" phrasing and suggests more human rewrites.
- **Regenerate with tweak** — "make it punchier / shorter / more casual" per output.

- **Save & reuse** multiple voice profiles (e.g., personal vs. brand).

## 5. Why AI is the core function

AI is not decorative — it *is* the product:

- **Voice extraction:** analyzing samples to build a style representation.
- **Platform-conditioned rewriting:** generating distinct, format-aware output while constraining to the captured voice.

This satisfies the challenge's two separate requirements: (a) IBM Bob as the primary build assistant, and (b) AI as a core user-facing feature of the finished app.

## 6. Tech stack

| Layer         | Choice                                              | Why                                         |
|---------------|-----------------------------------------------------|---------------------------------------------|
| Front end     | Next.js + React + Tailwind                          | Fast to build, great demo polish            |
| AI            | IBM Granite / watsonx (or a hosted LLM API)         | Recommended challenge tech; core generation |
| Orchestration | LangChain (optional)                                | Manage prompts / voice-profile context      |
| State/storage | Local storage or lightweight DB (SQLite / Supabase) | Save voice profiles; keep it simple         |
| Hosting       | Vercel                                              | One-click public deploy for judges          |




Keep it lean — a solo build with a 40-Bobcoin / 30-day trial. Favor a clean single-page workflow over feature sprawl.

## 7. Architecture (simple)

User samples → [Voice Profile Builder] → **voiceProfile** (JS ON style descriptor)

Pillar text ————— |



## 8. How IBM Bob will be used (document this for judges)

Use Bob meaningfully and keep notes/screenshots for the submission:

- Understanding requirements and planning the app (Plan mode)
- Scaffolding the Next.js project and components (Code mode)
- Designing the voice-profile data structure
- Writing and refining the generation prompts
- Debugging API calls and state handling
- Generating tests and improving the README

## 9. README outline (required, must cover the Creative Industries theme)

- Problem being solved (repurposing + voice loss)
- The solution and how it works
- **Selected challenge theme: Reimagine Creative Industries with AI**
- AI and technical approach
- Project architecture
- How IBM Bob was used
- Install / run instructions

- Link to live demo + demo video

## 10. Submission checklist

- ☐ Functioning prototype (core workflow works end-to-end)
- ☐ IBM Bob used as primary dev tool (documented)
- ☐ AI is a central, user-facing function
- ☐ IBM SkillsBuild certificate completed + saved
- ☐ Public GitHub repo with full README
- ☐ Public demo video ( $\leq 3$  minutes)
- ☐ Project Page published (not draft), all sections complete
- ☐ All links tested in incognito/private browser

## 11. Two-week build timeline (to July 31)

| Dates     | Focus                                                                                                                                |
|-----------|--------------------------------------------------------------------------------------------------------------------------------------|
| Jul 17–18 | Register on platform; start IBM Bob trial; complete a SkillsBuild activity + save certificate; scaffold Next.js repo (public GitHub) |
| Jul 19–21 | Build Voice Profile Builder + pillar input UI; wire up first platform (LinkedIn) generation                                          |
| Jul 22–24 | Add remaining platforms (X, Instagram, YouTube, newsletter); tune platform-specific prompts                                          |
| Jul 25–26 | Voice-match refinement; copy/export; basic styling polish                                                                            |
| Jul 27–28 | Deploy to Vercel; write full README; internal testing + fixes                                                                        |
| Jul 29    | Record $\leq 3$ -min demo video; publish Project Page                                                                                |
| Jul 30    | Buffer: retest all links in incognito, final polish, <b>submit early</b>                                                             |
| Jul 31    | Reserve only as emergency buffer (don't rely on the SAST deadline conversion)                                                        |

## 12. How it scores against the judging rubric

- **Technical Execution** — real working end-to-end flow, clear Bob usage.

- **Innovation** — voice-preserving repurposing is a fresh angle vs. generic "generate a caption" tools.
- **Challenge Fit** — directly improves a real creative workflow (content production).
- **Implementation & Feasibility** — genuinely usable by solo creators today; clear path to scale.

## 13. Risks & mitigations

- **Bobcoin / trial limits** → push code to GitHub constantly; keep a second email ready for a new trial if needed.
- **Voice match feels weak in demo** → pick a distinctive sample voice for the video so the effect is obvious.
- **Scope creep** → ship the 5-platform MVP first; treat authenticity check as bonus only.

## 14. Costs & AI hosting



The challenge does **not** cover your app's AI/API running costs. The IBM Bob free trial (40 Bobcoins / 30 days) only covers Bob as your *development* tool — not the AI running inside EchoDraft. Good news: this can be built and run for **R0 / \$0**.

### Two different AI requirements, two different funding sources

| Requirement                                 | Who pays                                        |
|---------------------------------------------|-------------------------------------------------|
| IBM Bob as primary dev tool                 | ✅ Covered by free trial (40 Bobcoins / 30 days) |
| AI as a core function <i>inside the app</i> | ❌ On me — the challenge does not cover this     |

### Full cost breakdown

| Item               | Cost             | Notes                                                          |
|--------------------|------------------|----------------------------------------------------------------|
| IBM Bob (dev tool) | Free             | 40 Bobcoins / 30-day trial                                     |
| App's AI / LLM API | Free             | Gemini Flash / Groq free tier, or local Granite                |
| Hosting            | Free             | Vercel free tier                                               |
| Domain             | Free             | Default <code>.vercel.app</code> URL — no custom domain needed |
| GitHub repo        | Free             | Public repo required anyway                                    |
| <b>Total</b>       | <b>~R0 / \$0</b> | Fully achievable at zero cost                                  |

## AI hosting options (cheapest first)

1. **Free LLM API (recommended for the demo)** — Google AI Studio (Gemini Flash) offers a genuinely free tier (~1.5M tokens/day, no credit card). Groq, OpenRouter, Mistral, and Cohere also have permanent free tiers. Easiest path: drop the API key in the app and go.
2. **Run IBM Granite locally (best "challenge fit" story)** — Granite models are open source (Apache 2.0) and free to run via LM Studio or Ollama. Great narrative for judges; the catch is the model runs on my machine, so the public demo needs it hosted or demoed locally in the video.
3. **Avoid watsonx.ai paid plans** — the Standard plan is ~\$1,050+/month. Use only the free trial, or stick with option 1 or 2.

## Recommended approach

Build and demo with **Gemini Flash's free API** (zero cost, no hardware needs, public demo works), and state in the README/video that EchoDraft is **designed to run on IBM Granite** (a swappable open model). This gives \$0 running cost, a working public demo, and a strong "uses IBM's recommended tech" story.



Watch out: free tiers rate-limit requests. Pre-test all demo examples before recording the 3-minute video so I don't hit a limit mid-recording.